

CARNEGIE MELLON UNIVERSITY  
DEPARTMENT OF COMPUTER SCIENCE  
15-445/645 – DATABASE SYSTEMS (FALL 2019)  
PROF. ANDY PAVLO

Homework 2 (by Amadou Ngom)  
Due: **Monday Sept 30, 2019 @ 11:59pm**

**IMPORTANT:**

- **Upload this PDF** with your answers to **Gradescope by 11:59pm on Monday Sept 30, 2019**.
- **Plagiarism:** Homework may be discussed with other students, but all homework is to be completed **individually**.

For your information:

- Graded out of **100** points; **4** questions total
- Rough time estimate:  $\approx$ 1-4 hours (0.5-1 hours for each question)

*Revision : 2019/09/25 15:06*

| Question           | Points | Score |
|--------------------|--------|-------|
| Cuckoo Hashing     | 20     |       |
| B+Tree             | 45     |       |
| Extendible Hashing | 25     |       |
| Suffix Trees       | 10     |       |
| Total:             | 100    |       |

---

**Number of Days this Assignment is Late:**

---

**Number of Late Day You Have Left:**

**Question 1: Cuckoo Hashing.....[20 points]**

Consider the following cuckoo hashing schema:

- Both tables have a size of 4.
- The hashing function of the first table returns the lowest two bits:  $h_1(x) = x \& 0b11$ .
- The hashing function of the second table returns the next two bits:  $h_2(x) = (x \gg 2) \& 0b11$ .
- When replacement is necessary, first select an element in the second table.
- The original content is shown in Figure 1.

|      | Table 1 | Table 2 |
|------|---------|---------|
| 00 → | 4       |         |
| 01 → |         |         |
| 10 → | 14      |         |
| 11 → |         |         |

Figure 1: Initial contents of the hash tables.


Use the following template to answer the questions: <https://cmudb.io/fall2019-hw1>.

- (a) [4 points] Insert keys 12 and 10. Draw the resulting two tables.


$12 = 8 + 4 = 1100$      $h_1(12) = 00$      $h_2(12) = 11$   
 $10 = 8 + 2 = 1010$      $h_1(10) = 10$      $h_2(10) = 10$

|    |    |    |
|----|----|----|
| 00 | 4  |    |
| 01 |    |    |
| 10 | 14 | 10 |
| 11 |    | 12 |

(b) [4 points] Then delete 14, and insert 8. Draw the resulting two tables.



(c) [6 points] Finally, insert 28. Draw the resulting two tables.

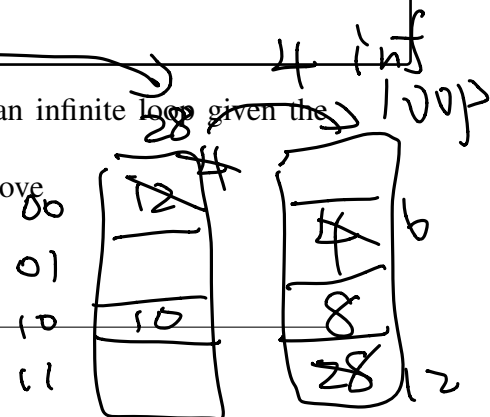


(d) [6 points] What is the smallest key that potentially causes an infinite loop given the tables in (c)

- ☐ 0   ☐ 2   ☐ 5   ☒ 6   ☐ 7   ☐ 9   ☐ None of the above


$0010$     $0011$     $0110$   
                        

Homework 2 continues...



**Question 2: B+Tree.....[45 points]**


Consider the following B+tree.

Figure 2: B+ Tree of order  $d = 4$  and height  $h = 2$ .

When answering the following questions, be sure to follow the procedures described in class and in your textbook. You can make the following assumptions:

- A left pointer in an internal node guides towards keys  $<$  than its corresponding key, while a right pointer guides towards keys  $\geq$ .
- A leaf node underflows when the number of **keys** goes below  $\lceil \frac{d-1}{2} \rceil$ .
- An internal node underflows when the number of **pointers** goes below  $\lceil \frac{d}{2} \rceil$ .



Use the following draw.io template for your answers:

<https://cmudb.io/fall2019-hw2>(a) **[15 points]** Insert  $10^*$  into the B+tree. Draw the resulting tree.(b) **[5 points]** How many pointers (parent-to-child and sibling-to-sibling) do you chase to find all keys between 5 and 15?
☐ 2   ☐ 3   ☒ 4   ☐ 5   ☐ 6   ☐ 7

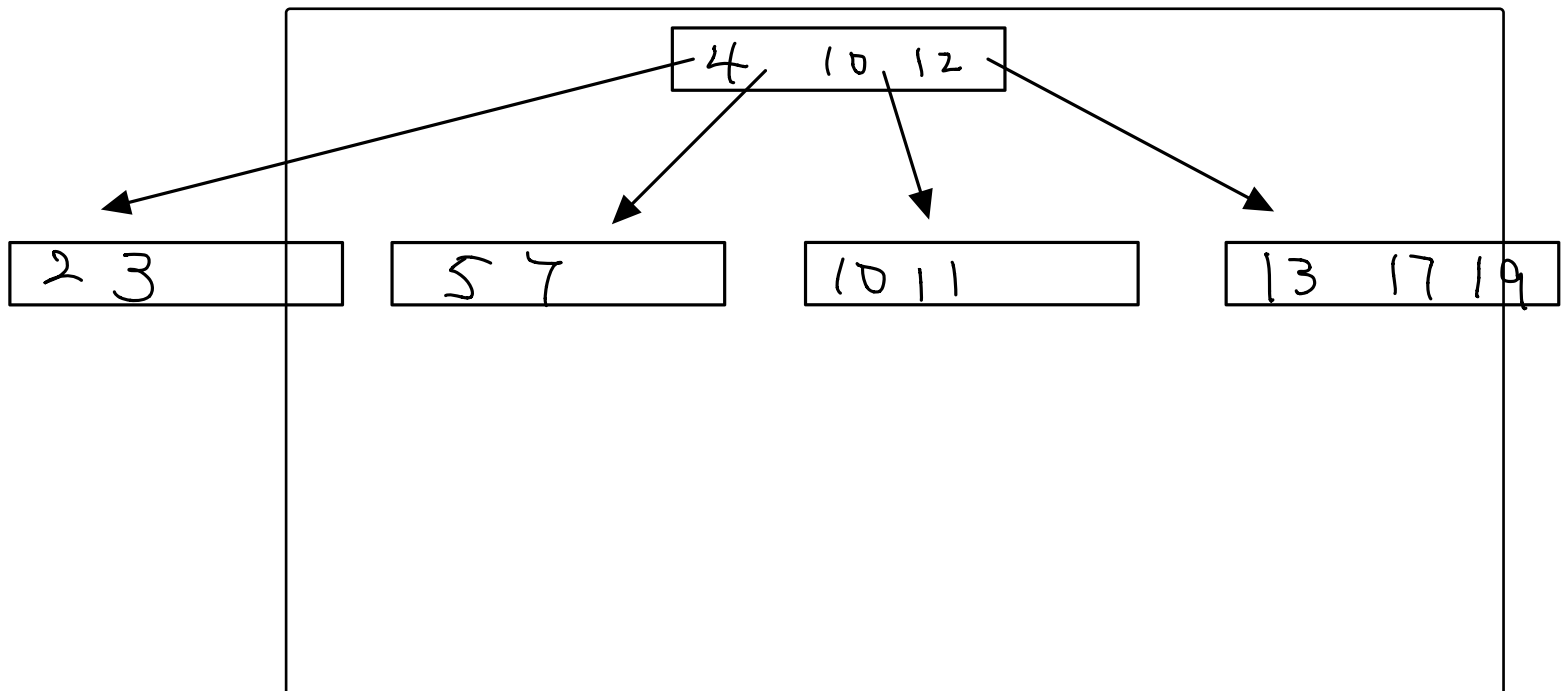
Find 5, then keep going right

Question 2 continues...


(a)




(c) [15 points] Then delete 23\*. Draw the resulting tree.



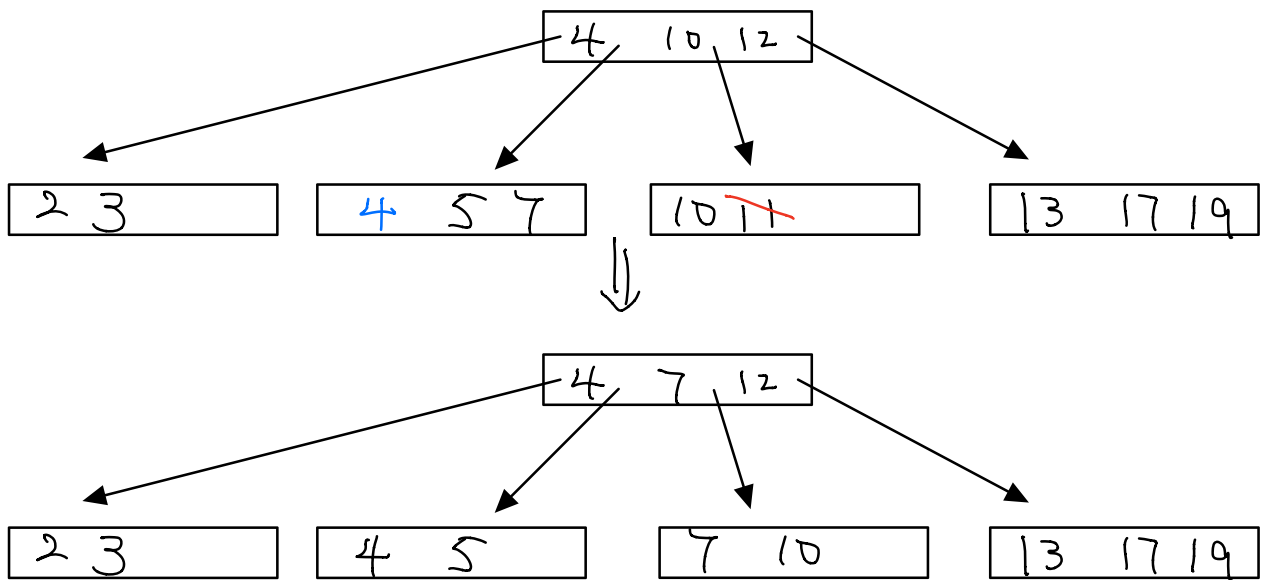
(d) [10 points] Finally insert 4\* and delete 11\*. Draw the resulting tree.



(b)



(c)



**Question 3: Extendible Hashing.....[25 points]**

Consider an extendible hashing structure such that:

- Each bucket can hold up to two records.
- The hashing function uses the lowest  $g$  bits, where  $g$  is the global depth.

(a) Starting from an empty table, insert keys 15, 3, 7, 14.

i. [3 points] What is the global depth of the resulting table?

- ☐ 0   ☐ 1   ☐ 2   ☒ 3   ☐ 4   ☐ None of the above

ii. [3 points] What is the local depth the bucket containing 14?

- ☐ 0   ☒ 1   ☐ 2   ☐ 3   ☐ 4   ☐ None of the above

iii. [3 points] What is the local depth of the bucket containing 3?

- ☐ 0   ☐ 1   ☐ 2   ☒ 3   ☐ 4   ☐ None of the above

(b) Starting from the result in (a), you insert keys 1, 9, 23, 11, 17.

i. [4 points] Which key will first cause a split (without doubling the size of the table)?


- ☐ 1   ☐ 9   ☐ 23   ☐ 11   ☒ 17   ☐ None of the above

ii. [4 points] Which key will first make the table double in size?

- ☐ 1   ☐ 9   ☒ 23   ☐ 11   ☐ 17   ☐ None of the above

(c) Now consider the table below, along with the following deletion rules:

1. If two buckets have the same local depth  $d$ , and share the first  $d - 1$  bits of their indexes (e.g. 010 and 110 share the first 2 bits), then they can be merged if the total capacity fits in a single bucket. The resulting local depth is  $d - 1$ .
2. If the global depth  $g$  becomes strictly greater than all local depths, then the table can be halved in size. The resulting global depth is  $g - 1$ .



Question 3 continues...



Starting from the table above, delete keys 2, 7, 13, 15, 29.


i. **[4 points]** Which deletion first causes a reduction in a local depth.

☐ 2   ☐ 7   ☒ 13   ☐ 15   ☐ 29   ☐ None of the above



ii. **[4 points]** Which deletion first causes a reduction in global depth.

☐ 2   ☐ 7   ☒ 13   ☐ 15   ☐ 29   ☐ None of the above

(a)  $g=0$



add 7



$$15 = 8 + 4 + 2 + 1$$

$$= 1111$$

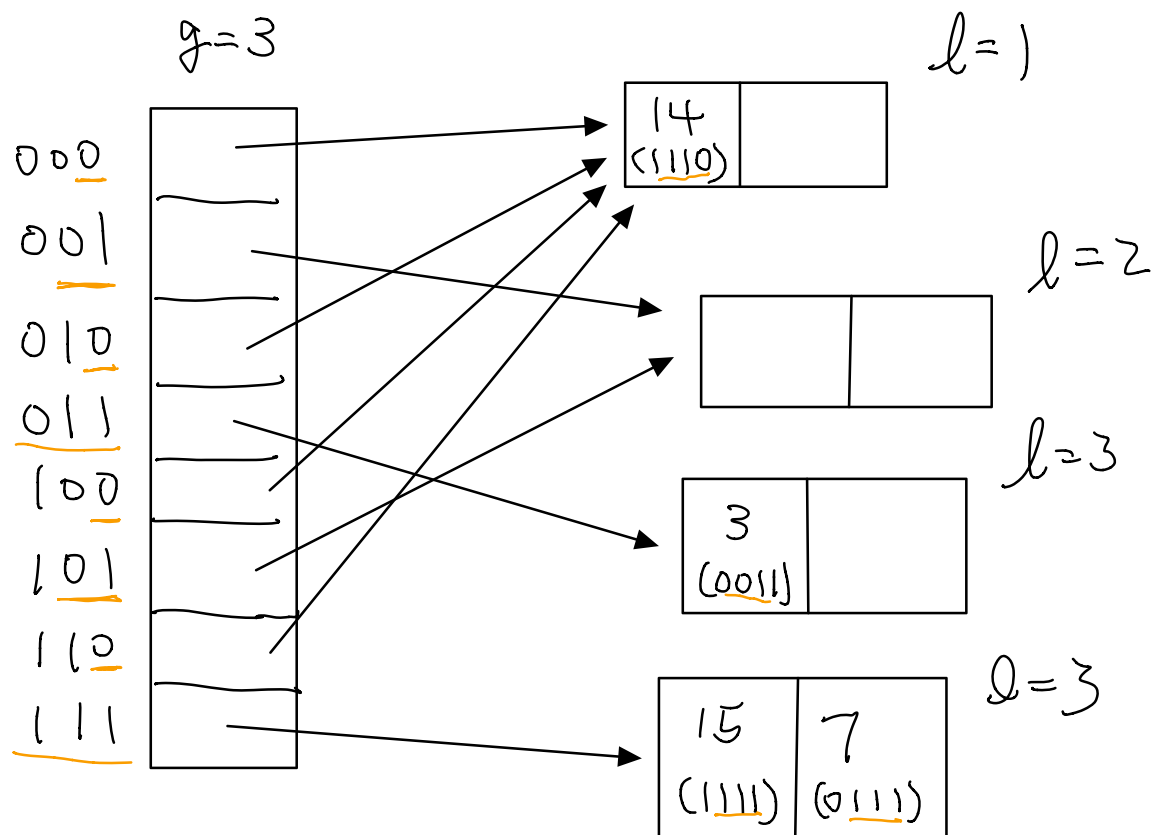
$$3 = 0011$$

$$7 = 0111$$


$$14 = 1110$$

$$7 = 0111$$

$\Rightarrow$  still full  
 $\Rightarrow$  split



(b)



1, 9, 23, 11, 17

1 = 00001

9 = 01001

$23 = 16 + 4 + 2 + 1 = 10111$

$11 = 8 + 2 + 1 = 01011$

$17 = 16 + 1 = 10001$

Split when  
23 is inserted

$g=4$

Insert 23

0000

0001

0010

0011

0100

0101

0110

0111

1000

1001

1010


1011

1100

1101

1110

1111



|        |  |
|--------|--|
| 14     |  |
| (1110) |  |

$l=1$

|         |         |
|---------|---------|
| 1       | 9       |
| (00001) | (01001) |

$l=2$

|         |  |
|---------|--|
| 3       |  |
| (00011) |  |

$l=3$

|         |         |
|---------|---------|
| 7       | 23      |
| (00111) | (10111) |

$l=4$

|         |  |
|---------|--|
| 15      |  |
| (01111) |  |

$l=4$

$g=4$

Insert 11 (01011)

$l=1$

0000

0001

0010

0011

0100

0101

0110

0111

1000

1001

1010

1011

1100

1101

1110

1111

|              |  |
|--------------|--|
| 14<br>(1110) |  |
|--------------|--|

$l=2$

|              |              |
|--------------|--------------|
| 1<br>(00001) | 9<br>(01001) |
|--------------|--------------|

$l=3$

|              |               |
|--------------|---------------|
| 3<br>(00011) | 11<br>(01011) |
|--------------|---------------|

$l=4$


|              |               |
|--------------|---------------|
| 7<br>(00111) | 23<br>(10111) |
|--------------|---------------|

$l=4$

|               |  |
|---------------|--|
| 15<br>(01111) |  |
|---------------|--|


$g=4$

Insert 17 (10001) (first part)  $l=1$



$g=4$

Insert 17(10001) (second part)  $l=1$





(C)

$g=3$

$l=1$

000

001

010


011

100


101

110


111




$l=3$




$l=2$




$l=3$






$\Rightarrow$  halve the size



**Question 4: Suffix Trees ..... [10 points]**

Consider the following suffix tree for **unsigned 32-bit integers**.



one hex  
= 4 bits

Figure 4: Suffix Tree

- (a) [3 points] Which of the following elements belong to the suffix tree. Select all that apply.

☐ 0x45BD0000   ☐ 0x0000CAAC   ☐ 0xFFAAAA00   ☐ 0xACCA0000   ☐ 0xBD000000  
☒ None of the above

- (b) [7 points] Insert the key 0x00FFAABB. Draw the resulting tree using this template: <https://cmudb.io/fall12019-hw4>.

